

#Launch virtual machines with 2gb ram, 2 cpu minimum for master node and 1gb ram, 2 cpu for slave nodes in virtual box. I am using Ubuntu 20.04.5 LTS

Install docker on both master and slave vm:

```
sudo apt-get update
```

```
sudo apt-get install ca-certificates curl gnupg lsb-release
```

```
sudo mkdir -m 0755 -p /etc/apt/keyrings
```

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o  
/etc/apt/keyrings/docker.gpg
```

```
echo "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg]  
https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" | sudo tee  
/etc/apt/sources.list.d/docker.list > /dev/null
```

```
sudo apt-get update
```

```
sudo apt-get install docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-  
plugin
```

#You can confirm docker installed with “docker version”

Install cri-dockerd on both master and slave vm. This adapter provides a shim for Docker Engine that lets you control Docker via the Kubernetes Container Runtime Interface.

```
git clone https://github.com/Mirantis/cri-dockerd.git
```

###Install GO###

```
wget https://storage.googleapis.com/golang/getgo/installer_linux
```

```
sudo chmod +x ./installer_linux
```

```
./installer_linux
```

```
source ~/.bash_profile
```

You can confirm go installed with “go version” command

```
cd cri-dockerd
```

```
mkdir bin
```

```
go build -o bin/cri-dockerd #This build will take some time
```

```
mkdir -p /usr/local/bin
```

```
sudo install -o root -g root -m 0755 bin/cri-dockerd /usr/local/bin/cri-dockerd
```

```
sudo cp -a packaging/systemd/* /etc/systemd/system
```

```
sudo sed -i -e 's,/usr/bin/cri-dockerd,/usr/local/bin/cri-dockerd,' /etc/systemd/system/cri-  
docker.service
```

```
sudo systemctl daemon-reload  
sudo systemctl enable cri-docker.service  
sudo systemctl enable --now cri-docker.socket
```

Installing kubeadm, kubelet and kubect1 (on both master and slave)

```
sudo apt-get update  
sudo apt-get install -y apt-transport-https ca-certificates curl
```

```
sudo curl -fsSLo /etc/apt/keyrings/kubernetes-archive-keyring.gpg  
https://packages.cloud.google.com/apt/doc/apt-key.gpg
```

```
echo "deb [signed-by=/etc/apt/keyrings/kubernetes-archive-keyring.gpg]  
https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee  
/etc/apt/sources.list.d/kubernetes.list
```

```
sudo apt-get update  
sudo apt-get install -y kubelet kubeadm kubect1  
sudo apt-mark hold kubelet kubeadm kubect1
```

Forwarding IPv4 and letting iptables see bridged traffic (only on Master vm)

```
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf  
overlay  
br_netfilter  
EOF
```

```
sudo modprobe overlay  
sudo modprobe br_netfilter
```

sysctl params required by setup, params persist across reboots

```
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf  
net.bridge.bridge-nf-call-iptables = 1  
net.bridge.bridge-nf-call-ip6tables = 1  
net.ipv4.ip_forward = 1  
EOF
```

Apply sysctl params without reboot

```
sudo sysctl --system
```

#Run the below command only on master node (192.168.29.46 is master node ip, please use your master vm ip. Also use same pod network cidr given below):

```
sudo kubeadm init --apiserver-advertise-address=192.168.29.46 --pod-network-cidr=10.244.0.0/16 --  
cri-socket=unix:///var/run/cri-dockerd.sock
```

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

A token with join command will be displayed, copy and paste that in slave node to join the cluster. At the end "--cri-socket= unix:///var/run/cri-dockerd.sock" is added as I was getting Found multiple CRI endpoints on the host error

```
sudo kubeadm join 192.168.29.46:6443 --token 8te8jc.dtic56eqnwufo3w --discovery-token-ca-cert-hash sha256:6fff1209a717cb7d1928e1027827d9a7e186c7fb4f161044d9dfac67ce9000f5 --cri-socket=unix:///var/run/cri-dockerd.sock (# run the token generated by your master vm & wait for it to complete)
```

Run the below command only on master node:

```
wget https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml
```

```
vi kube-flannel.yml
```

```
containers:
- args:
  - --ip-masq
  - --kube-subnet-mgr
  - --iface=enp0s8
```

In args add --iface=enp0s8 , here enp0s8 is master node ethernet adaptor. You can check your master node ethernet adaptor with ifconfig/ip add show command. After that save this file :wq

```
kubectl apply -f kube-flannel.yml (# run on master node)
```

#Master node output:

```
MINGW64~/e/Projects/kubeadm/kubemaster
vagrant@kubemaster:~$ kubectl get nodes
NAME           STATUS    ROLES    AGE   VERSION
kubemaster     Ready     control-plane   8m53s   v1.26.1
kubenode1      Ready     <none>         2m1s    v1.26.1
kubenode2      Ready     <none>         87s     v1.26.1
vagrant@kubemaster:~$ kubectl get pods -A
NAMESPACE     NAME                                     READY   STATUS    RESTARTS   AGE
kube-flannel   kube-flannel-ds-dnf6m                  1/1     Running   0           4m45s
kube-flannel   kube-flannel-ds-gfdhr                  1/1     Running   0           90s
kube-flannel   kube-flannel-ds-k7pzk                  1/1     Running   0           2m4s
kube-system    coredns-787d4945fb-25krj               1/1     Running   0           8m7s
kube-system    coredns-787d4945fb-fcf7p               1/1     Running   0           8m7s
kube-system    etcd-kubemaster                        1/1     Running   0           8m33s
kube-system    kube-apiserver-kubemaster               1/1     Running   0           8m33s
kube-system    kube-controller-manager-kubemaster      1/1     Running   1 (3m45s ago)  8m33s
kube-system    kube-proxy-gkfxb                        1/1     Running   0           8m7s
kube-system    kube-proxy-rxbh2                        1/1     Running   0           2m4s
kube-system    kube-proxy-vxmpx                        1/1     Running   0           90s
kube-system    kube-scheduler-kubemaster               1/1     Running   2 (96s ago)   8m33s
vagrant@kubemaster:~$
```